

Style Transfer

An Introduction

Eka Kurniawan

BukaLapak

Background


Style Transfer



Bukalapak



Introduction



VGG19 Model

Invert Representation

This section introduces our method to compute an approximate inverse of an image representation. This is formulated as the problem of finding an image whose representation best matches the one given [34]. Formally, given a representation function $\Phi : \mathbb{R}^{H \times W \times C} \rightarrow \mathbb{R}^d$ and a representation $\Phi_0 = \Phi(\mathbf{x}_0)$ to be inverted, reconstruction finds the image $\mathbf{x} \in \mathbb{R}^{H \times W \times C}$ that minimizes the objective:

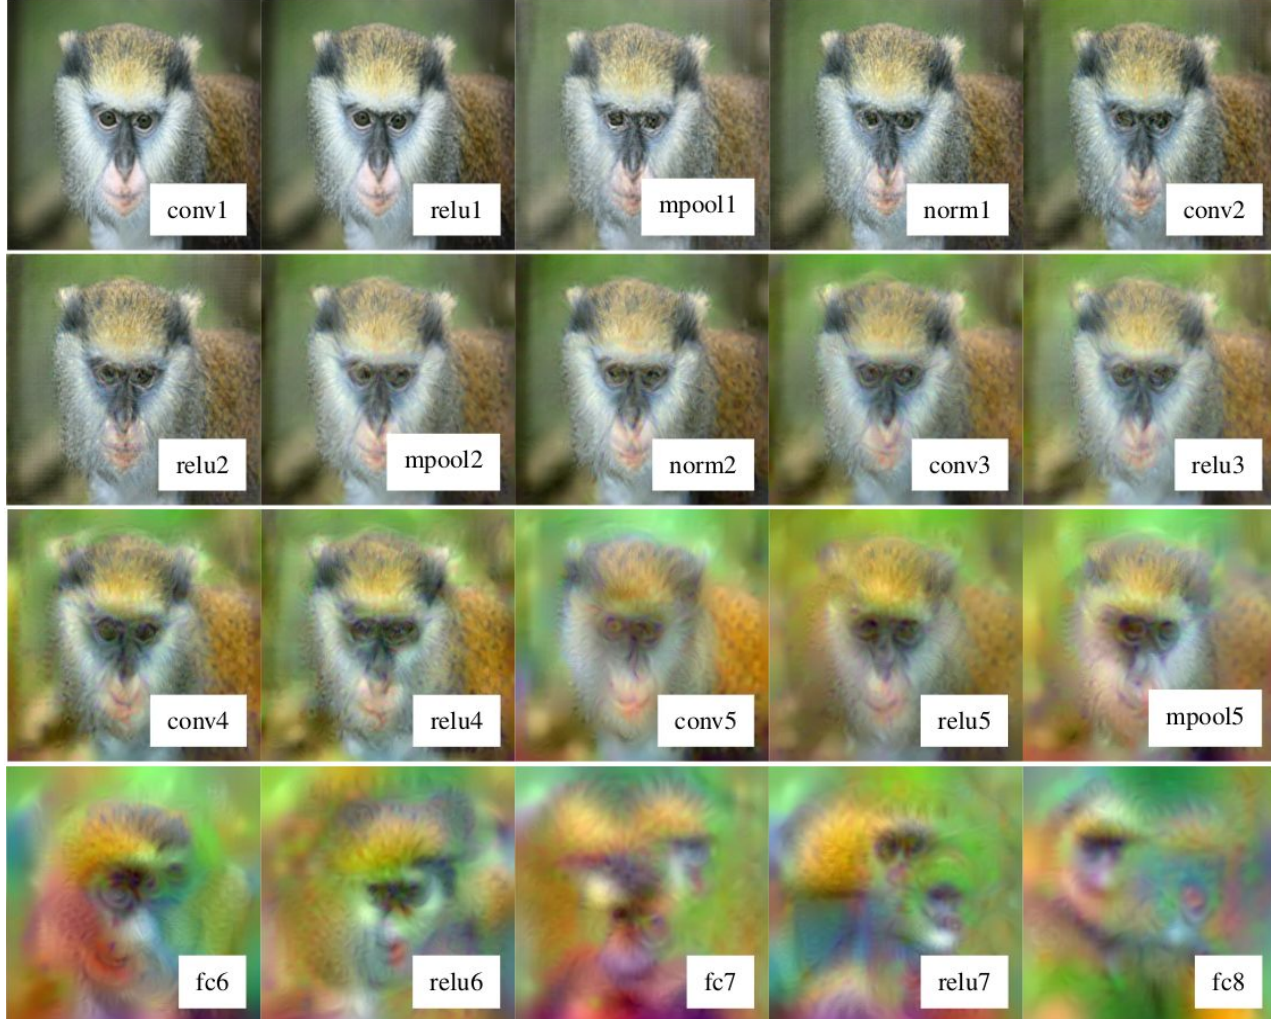
$$\mathbf{x}^* = \underset{\mathbf{x} \in \mathbb{R}^{H \times W \times C}}{\operatorname{argmin}} \ell(\Phi(\mathbf{x}), \Phi_0) + \lambda \mathcal{R}(\mathbf{x}) \quad (1)$$


where the loss ℓ compares the image representation $\Phi(\mathbf{x})$ to the target one Φ_0 and $\mathcal{R} : \mathbb{R}^{H \times W \times C} \rightarrow \mathbb{R}$ is a regulariser capturing a *natural image prior*.



Invert Representation

Invert Representation



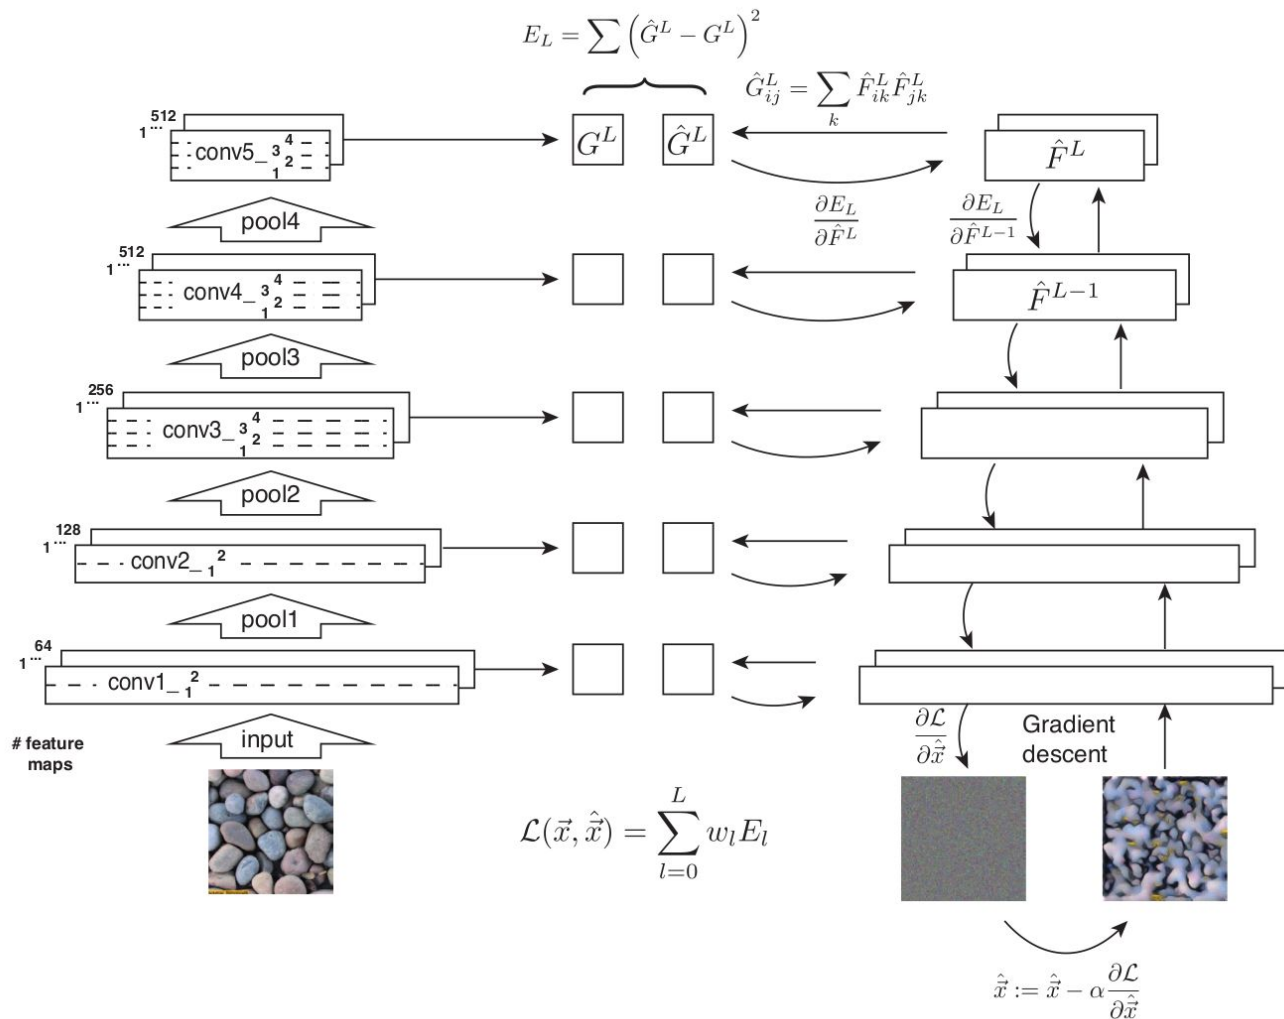


Gram Matrix


Localization Removal

to extract just the style information from the features without the information where the style is located in the image


Texture Synthesis



Texture Synthesis




conv1_1




pool1




pool2



pool3




pool4




original

Texture Synthesis



conv1_1



pool1



pool2



pool3



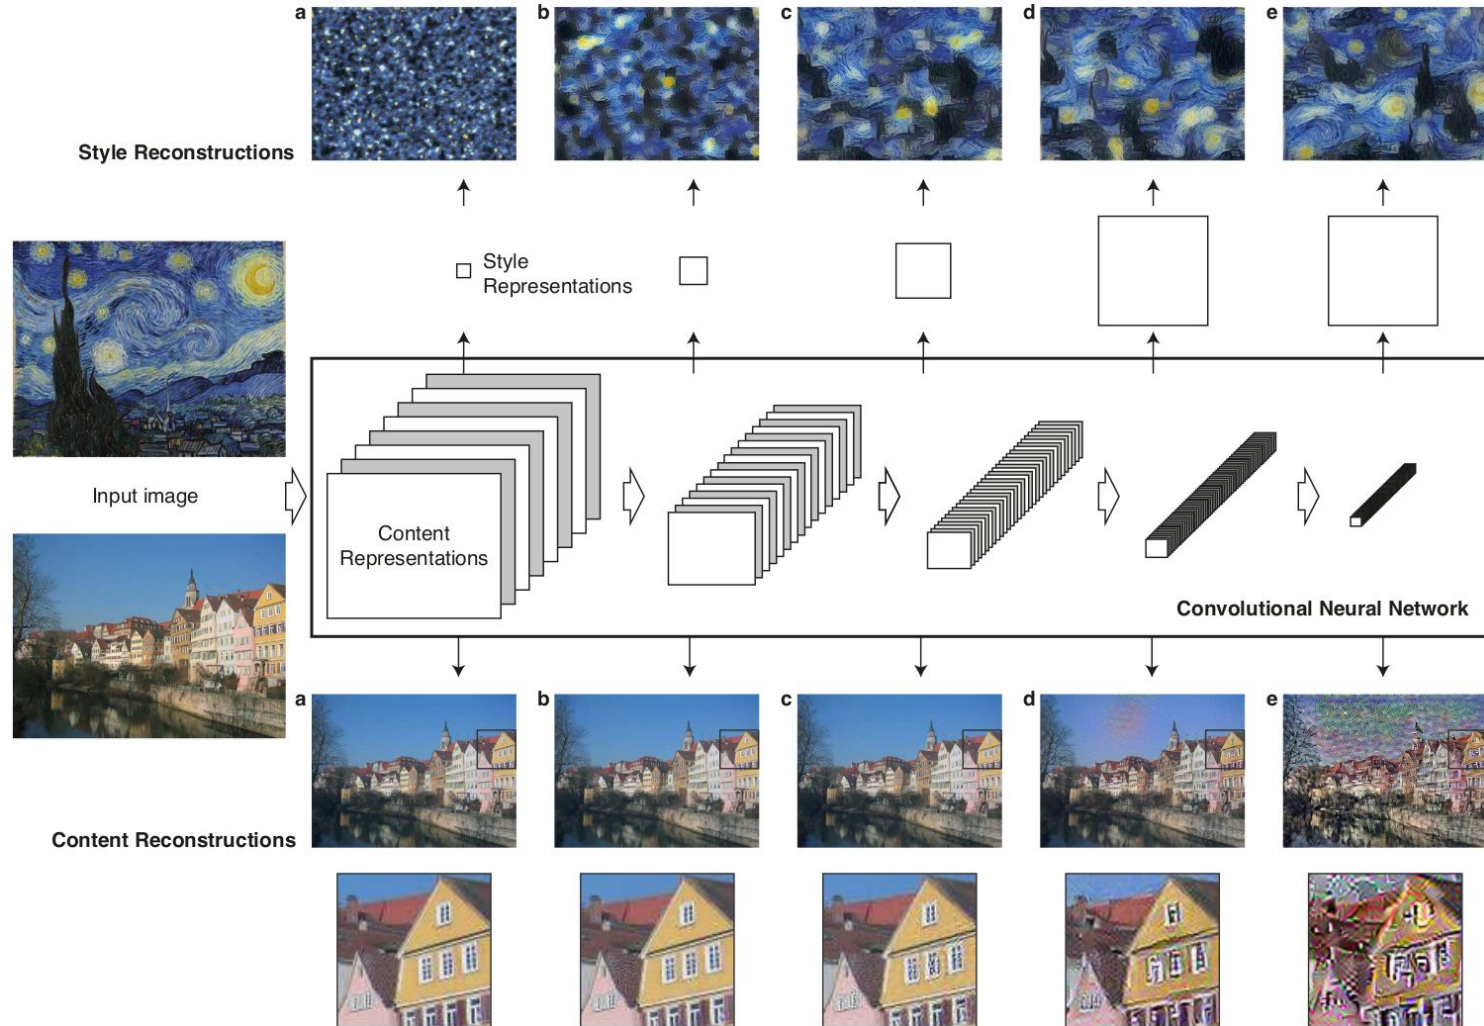
pool4

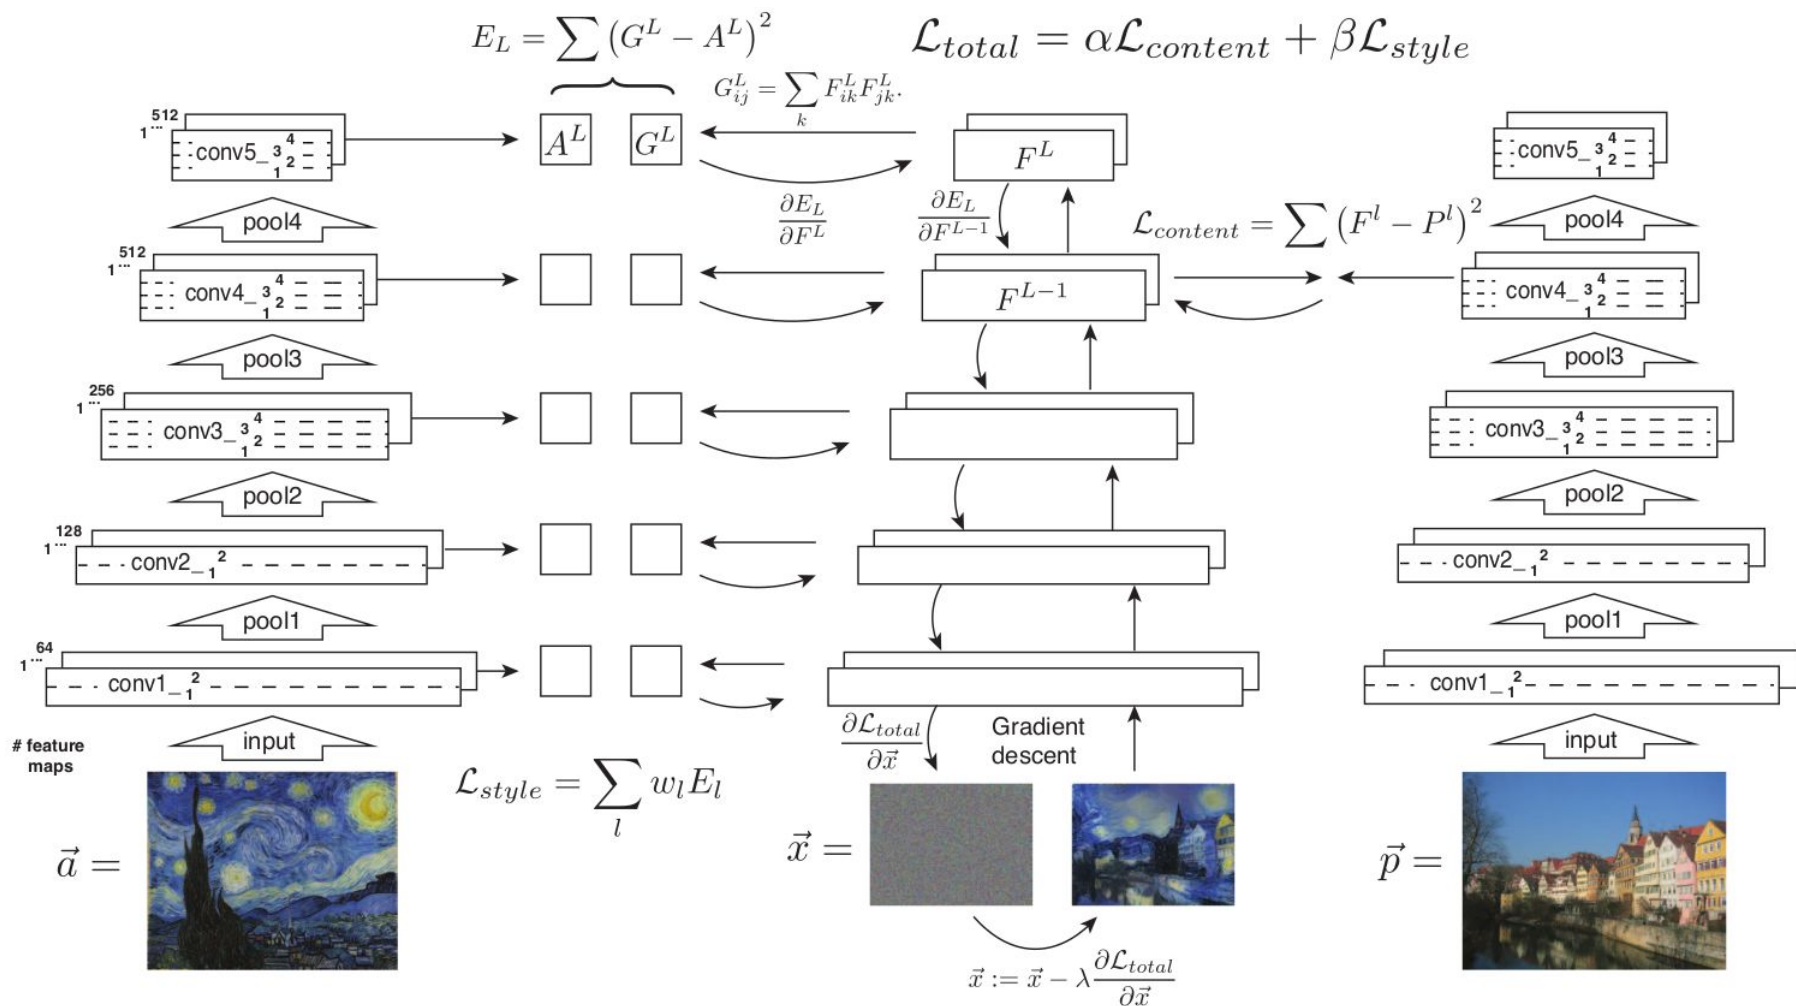


original


Style Transfer

Invert Representation + Texture Synthesis (Style and Content Reconstructions)





Implementation



VGG19 Model

Style

`conv1_1`
`conv2_1`
`conv3_1`
`conv4_1`
`conv5_1`

Content

`conv4_2`


```
# get the "features" portion of VGG19 (we will not need the "classifier" portion)  
vgg = models.vgg19(pretrained=True).features  
  
# freeze all VGG parameters since we're only optimizing the target image  
for param in vgg.parameters():  
    param.requires_grad_(False)
```



```
(0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(1): ReLU(inplace)
(2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(3): ReLU(inplace)
(4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(6): ReLU(inplace)
(7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(8): ReLU(inplace)
(9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(11): ReLU(inplace)
(12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(13): ReLU(inplace)
(14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(15): ReLU(inplace)
(16): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(17): ReLU(inplace)
(18): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(19): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(20): ReLU(inplace)
(21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(22): ReLU(inplace)
(23): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(24): ReLU(inplace)
(25): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(26): ReLU(inplace)
(27): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(29): ReLU(inplace)
(30): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(31): ReLU(inplace)
(32): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(33): ReLU(inplace)
(34): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(35): ReLU(inplace)
(36): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
```

```
def get_features(image, model, layers=None):  
    """ Run an image forward through a model and get the features for  
        a set of layers. Default layers are for VGGNet matching Gatys et al (2016)  
    """  
  
    ## TODO: Complete mapping layer names of PyTorch's VGGNet to names from the paper  
    ## Need the layers for the content and style representations of an image  
    if layers is None:  
        layers = {'0': 'conv1_1', # for style  
                  '5': 'conv2_1', # for style  
                  '10': 'conv3_1', # for style  
                  '19': 'conv4_1', # for style  
                  '21': 'conv4_2', # for content  
                  '28': 'conv5_1', # for style  
                  }  
  
    ## -- do not need to change the code below this line -- ##  
    features = {}  
    x = image  
    # model._modules is a dictionary holding each module in the model  
    for name, layer in model._modules.items():  
        x = layer(x)  
        if name in layers:  
            features[layers[name]] = x  
  
    return features
```

```
def gram_matrix(tensor):  
    """ Calculate the Gram Matrix of a given tensor  
        Gram Matrix: https://en.wikipedia.org/wiki/Gramian\_matrix  
    """  
  
    ## get the batch_size, depth, height, and width of the Tensor  
    ## reshape it, so we're multiplying the features for each channel  
    ## calculate the gram matrix  
    tensor_shape = tensor.shape  
    tensor_flat = tensor.view(tensor_shape[0], tensor_shape[1], -1)  
    gram = torch.bmm(tensor_flat, torch.transpose(tensor_flat, 1, 2))  
  
    return gram
```



```
# get content and style features only once before forming the target image
content_features = get_features(content, vgg)
style_features = get_features(style, vgg)

# calculate the gram matrices for each layer of our style representation
style_grams = {layer: gram_matrix(style_features[layer]) for layer in style_features}

# create a third "target" image and prep it for change
# it is a good idea to start of with the target as a copy of our *content* image
# then iteratively change its style
target = content.clone().requires_grad_(True).to(device)
```

```
# weights for each style layer
# weighting earlier layers more will result in *larger* style artifacts
# notice we are excluding `conv4_2` our content representation
style_weights = {'conv1_1': 0.9, #bar 1.,
                 'conv2_1': 0.8, #bar 0.8,
                 'conv3_1': 0.5, #bar 0.5,
                 'conv4_1': 0.3, #bar 0.3,
                 'conv5_1': 0.1} #bar 0.1}

# you may choose to leave these as is
content_weight = 1 # alpha
style_weight = 1e6 #bar 1e6 # beta
```



```
# iteration hyperparameters  
optimizer = optim.Adam([target], lr=0.003)  
steps = 1000 #bar 2000 # decide how many iterations to update your image (5000)
```

```
for ii in range(1, steps+1):
    target_features = get_features(target, vgg)
    content_loss = torch.mean((target_features['conv4_2'] - content_features['conv4_2']) ** 2)

    style_loss = 0
    # iterate through each style layer and add to the style loss
    for layer in style_weights:
        # get the "target" style representation for the layer
        target_feature = target_features[layer]
        _, d, h, w = target_feature.shape

        target_gram = gram_matrix(target_feature)

        style_gram = style_grams[layer]
        layer_style_loss = style_weights[layer] * \
            torch.mean((target_gram - style_gram) ** 2)

        # add to the style loss
        style_loss += layer_style_loss / (d * h * w)

    total_loss = (content_weight * content_loss) + \
        (style_weight * style_loss)


    optimizer.zero_grad()
    total_loss.backward()
    optimizer.step()
```



style



content



Conclusions

01 Features that represent style of an input image can be acquired from layers at the beginning of VGG19 model architecture whereas content is from the end.

02 Gram matrix can be used to remove localization.

03 Fully connected layers also lost the localization.

04 Style transfer optimizes features instead of parameters.

Thank You



Eka Kurniawan

AI Engineer

